

Lab 1

Part 1

I followed [this](#) tutorial on how to write and insert a basic module into the kernel. I first installed the required packages needed for kernel development. I then learned that you need to write your kernel init and exit functions and tell the kernel which functions to call when loading or unloading the kernel using the `module_init` and `module_exit` functions. The init and exit functions simply used `printk()` to print "Hello World" and "Goodbye World" respectively. The most difficult part was getting the module to compile, as I didn't realize that it wouldn't compile unless module documentation is added. After I compiled and successfully inserted the module into the kernel, the kernel reported that it was "tainted". Upon further research I found that it is just a warning message saying non-vendor approved software had been inserted into the kernel module.

Source code used in part 1:

I used the following source code from [Rober Oliver's "Writing a Simple Kernel Module"](#) tutorial:

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Robert W. Oliver II");
MODULE_DESCRIPTION("A simple example Linux module.");
MODULE_VERSION("0.01");

static int __init lkm_example_init(void) {
    printk(KERN_INFO "Hello, World!\n");
    return 0;
}

static void __exit lkm_example_exit(void) {
    printk(KERN_INFO "Goodbye, World!\n");
}

module_init(lkm_example_init);
module_exit(lkm_example_exit);
```

I changed the name of the init and exit functions from `lkm_example_init` and `lkm_example_exit` to `helloWorld` and `goodbyeWorld`. I also changed the module description, author, and version.

I also used the makefile from the same tutorial:

```
obj-m += lkm_example.o

all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

I only changed the name of the module from `lkm_example` to `hello`.

Part 2

At first, I thought we were supposed to create a `sysfs` file at the location of `/sys/module/<module_name>/parameters/<parameter_name>` using `k_objects`. Once I realized that we were just supposed to use module parameters, the problem became simpler. To implement the enable logging parameter, I followed another online tutorial to learn how to use the `module_param` macro. I created an integer parameter called `enable_logging` and made each `printk()` in the module code statement dependent on the value of `enable_logging`. After this, I ran into a permission issue. The `sudo echo 0 > /sys/module/hello/parameters/enable_logging` was not able to actually edit the parameter. I tried setting the write and read privileges for the module to global, but the module wouldn't compile. I then learned that you are not allowed to set the write privileges to a kernel module to global. I then found that the "sudo" in the previous command was only "working" on the echo and the redirection did not get sudo privileges. I then used `sudo su` to act as the root and set the module parameters to user read and write, and the tests worked as expected.

For the `double_me` parameter, I modified some code from an online tutorial on how to communicate with device drivers. The most difficult part of this assignment was finding the correct macro—`module_param_cb`—and online tutorial to use as `moduleparam.h` was difficult to understand. From the tutorial, I learned that `module_param_cb` macro takes an argument of a `kernel_param_ops` struct. This tells the macro what methods to call when the parameter is set or accessed by the user. So, I created a method that would be called when the parameter was set by the user, which checked if the parameter value given by the user was an integer—using the macro `param_set_int`—and then doubled the value set. This method was then put in the `kernel_param_ops` struct as the `.set` method and the standard `param_get_int` macro was used for the `.get` method in the struct. The rest of the set up for the `double_me` module parameter was similar to the `enable_logging` parameter.

Source code used in part 2:

For the `enable logging` parameter, I used code from Lirhan B.H.'s [blog post](#) on kernel module parameters. I used and modified the following two lines from the post:

```
static int irq=10;  
module_param(irq,int,0660);
```

I changed the name from `irq` to `enable_logging` and changed the permissions from `0660` to `S_IRUGO | S_IRUSR`.

For the double_me parameter I used and modified some of the following code from an [EmbeTronicX tutorial](#) on passing arguments to device drivers.

```
int notify_param(const char *val, const struct kernel_param *kp)
{
    int res = param_set_int(val, kp); // Use helper for write variable
    if(res==0) {
        printk(KERN_INFO "Call back function called...\n");
        printk(KERN_INFO "New value of cb_valueETX = %d\n", cb_valueETX);
        return 0;
    }
    return -1;
}

const struct kernel_param_ops my_param_ops =
{
    .set = &notify_param, // Use our setter ...
    .get = &param_get_int, // .. and standard getter
};

module_param_cb(cb_valueETX, &my_param_ops, &cb_valueETX, S_IRUGO|S_IWUSR );
```

I changed the name of the parameter, kernel_param_ops struct, and the .set method to double_me, double_me_ops, and double_val respectively. In my version of the notify_param method, I doubled the double_me variable and printed a message to the kernel log instead. I also set my permissions to user read and write instead of global read and user write. Lastly, I returned the recommended EINVAL instead of -1 if the input given by the user was not an integer.

My part 2 code is also a modified version of my part 1 code, and sources for my part 1 code can be found in the part 1 section of this lab report.

Part 3

Looking back on the work needed for parts 1-5 of this lab, part 3 was definitely the most challenging. The first challenge was figuring out how to create the device file.

[Chapter 3](#) of the Linux Device Drivers textbook gave a good overview of how to allocate a character drive major/minor numbers and how to handle the character driver struct, as did many other sources. However, figuring out how to make a device file in the code without calling something like `mknod` in the shell was difficult. However, after some googling, I found an [EmbeTronicX tutorial](#) on how to create a device class and a device, so that the device file would show up when the module was inserted.

The next challenge was writing the methods for the `file_operations` struct: `llseek`, `read`, `write`, `ioctl`, `open`, `close`. The `ioctl` method and `open` method were probably the two most difficult methods to code. `Open` was deceptively simple. Since the `open` system call returns a file descriptor, I thought that the `open` method I was writing was supposed to return a file descriptor. Only after some more research did I realize that the `open` method is called to notify the drive that the device file has been opened and that the kernel handles file descriptors. After realizing that, I modeled my `open`, `close`, `read`, `write`, and `llseek` methods off of the ones described in Chapter 3 of the Linux Device Driver that the book was using for its example memory management driver. My biggest debugging challenge with these four methods was with the `read` and `write` method. I assumed that `copy_from/to_user` returned the amount of bytes copied, when in actuality it just returns 0 on success.

The `ioctl` method was definitely the most challenging one to write. I had two big issues while writing the `ioctl` methods. First, I didn't know I needed to define the `ioctl` commands using `#define` and the specified syntax. Second, I didn't realize that the unsigned long `arg` argument in the `ioctl` command was from user space, thus if I gave `ioctl` a user space pointer, I had to use `copy_from_user` to get its real value. Once I figured out these two bugs, fixing other minor problems became a lot easier. I found these well explained in another [EmbeTronicX tutorial](#) about `ioctl`, which was very helpful.

In implementing the sysfs files, the allocated parameter was quite similar to module parameters we had implemented in the past, so it was not difficult to use the `module_param` macro to define an integer parameter with read only permissions. The regions parameter was a bit more difficult since I had to figure out how to return a string version of my linked list of allocations when the user requested it. After I realized that you can use `sprintf` in kernel code—which I read in [this tutorial](#) on sysfs—I was able to write a method to traverse my linked list of allocated regions and return the relevant information to the user in string form.

Lastly, there were a few issues I faced when compiling my kernel module and running my test program. Similar to the issue the student faced in the Ed discussion post, I also struggled with the fact that when compiling a module spread over multiple files, the module name can not be the same as the name of any of the source files. Additionally, I also didn't realize for a while that the test program needs to be run with `sudo` privileges, or else it doesn't have permission to open any of the driver files.

Data:

The values below were averages over 15 trials. Data is in seconds.

Time it takes to write to every byte in a 512KB region: 0.386421s

Time it takes to read every byte in a 512KB region: 0.370846s

Source code for part 3:

For the main module init and exit functions found in "MemManager.c", I used code from an [EmbeTronicX tutorial](#) on the creation of device files: The source code is as follows:

The init function:

```
81 static int __init etx_driver_init(void)
82 {
83     /*Allocating Major number*/
84     if((alloc_chrdev_region(&dev, 0, 1, "etx_Dev")) < 0){
85         pr_err("Cannot allocate major number\n");
86         return -1;
87     }
88     pr_info("Major = %d Minor = %d \n",MAJOR(dev), MINOR(dev));
89
90     /*Creating cdev structure*/
91     cdev_init(&etx_cdev,&fops);
92
93     /*Adding character device to the system*/
94     if((cdev_add(&etx_cdev,dev,1)) < 0){
95         pr_err("Cannot add the device to the system\n");
96         goto r_class;
97     }
98
99     /*Creating struct class*/
100    if((dev_class = class_create(THIS_MODULE,"etx_class")) == NULL){
101        pr_err("Cannot create the struct class\n");
102        goto r_class;
103    }
104
105    /*Creating device*/
106    if((device_create(dev_class,NULL,dev,NULL,"etx_device")) == NULL){
107        pr_err("Cannot create the Device 1\n");
108        goto r_device;
109    }
110    pr_info("Device Driver Insert...Done!!!\n");
111    return 0;
112
113    r_device:
114        class_destroy(dev_class);
115    r_class:
116        unregister_chrdev_region(dev,1);
117        return -1;
118 }
```

The main change I made was removing the gotos and replacing them with the relevant lines of code that were going to be called anyway to increase readability. I also changed variable names to be more relevant in the context of the assignment. Additionally, I added these three lines of code from [chapter 3 of the Linux Device Driver](#) textbook:

```
cdev_init(&dev->cdev, &scull_fops);
dev->cdev.owner = THIS_MODULE;
dev->cdev.ops = &scull_fops;
```

I again changed variable names to be more relevant to the project.

The exit function:

```
/*
** Module exit function
*/
static void __exit etx_driver_exit(void)
{
    device_destroy(dev_class,dev);
    class_destroy(dev_class);
    cdev_del(&etx_cdev);
    unregister_chrdev_region(dev, 1);
    pr_info("Device Driver Remove...Done!!!\n");
}
```

I changed variable names to fit the context of the assignment and removed the print statement

For the open method, I modeled my open method off of the scull open method from Chapter 3 of the Linux Device Driver textbook (linked above):

```
int scull_open(struct inode *inode, struct file *filp)
{
    struct scull_dev *dev; /* device information */

    dev = container_of(inode->i_cdev, struct scull_dev, cdev);
    filp->private_data = dev; /* for other methods */

    /* now trim to 0 the length of the device if open was write
    if ( (filp->f_flags & O_ACCMODE) == O_WRONLY) {
        scull_trim(dev); /* ignore errors */
    }
    return 0;          /* success */
}
```

I changed variable names and removed the trim to length zero part of this method as it was unnecessary for this assignment.

I also modeled my `llseek` method after the `scull_llseek` implementation in the linux device driver textbook ([chapter 6](#)).

```
loff_t scull_llseek(struct file *filp, loff_t off, int whence)
{
    struct scull_dev *dev = filp->private_data;
    loff_t newpos;

    switch(whence) {
        case 0: /* SEEK_SET */
            newpos = off;
            break;

        case 1: /* SEEK_CUR */
            newpos = filp->f_pos + off;
            break;

        case 2: /* SEEK_END */
            newpos = dev->size + off;
            break;

        default: /* can't happen */
            return -EINVAL;
    }
    if (newpos < 0) return -EINVAL;
    filp->f_pos = newpos;
    return newpos;
}
```

I also changed variable names to fix the context of this assignment. I also added an if statement to check if the new offset was less than the size of the allocated region.

For the module parameters, I used the source code from Part 2 of this lab. I changed variable names to be more relevant to this assignment. Additionally, for the "regions"

parameter, I wrote a custom get method instead of a custom set method. See the "source code for part 2" for the citation of that code.

Part 4

The first step to creating the smart module was to use the `EXPORT_SYMBOLS()` macro to make functions from my part 3 module accessible to my new module. I had a bit of difficulty here because I didn't realize at first that I had to declare the functions I was using from the simple module in the header file used in the smart module before I could correctly use them. Additionally, I had some significant difficulty compiling the smart module code, before I realized that the easiest way to do that was to compile both modules together. Apparently, using a makefile to compile files from two different folders is complex and not the best idea. For that reason, I put a copy of my part 3 code in my part 4 folder for easy compilation.

My implementation of the init and exit functions from the smart module was almost the same as my init and exit function for part 3, so that wasn't too difficult. I did have to change a few names and get rid of references to the module parameter and `kobject`, as they were handled in the part 3 module. Finally, the new read and write methods (`smart_read` and `smart_write`), had a few small changes to the original read and write methods from part 3. Instead of checking if each request was for one byte, the smart methods now check if the amount of bytes requested plus the current offset is outside of the current region. Additionally, the `copy_to_user` and `copy_from_user` methods now depend on `count` for the number of bytes to transfer.

Data from Part 4:

The data for this section was collected in bytes and milliseconds. All values shown below were averaged over 15 trials

Table 1. Average values of time to read/write to every byte in the 512KB region.

Amount of Data transferred per call	Time to Read (ms)	Time to Write (ms)
1B	390.549713	391.501495
64B	7.2412	6.776801
1KB	0.593933	0.607667
64KB	0.0326	0.034667
512KB	0.034667	0.034667




Figure 1. Bytes read per read function call vs. time take to read 512KB




Figure 2. Bytes written per write function call vs. time take to write to 512KB

Analysis:

There is definitely a difference in throughput between reading/writing smaller amounts of bytes at a time vs larger amounts. In fact, there is such a large difference the data points can not be very nicely displayed in graph form. For example, when writing one byte at a time, you can write about $1.3 \cdot 10^6$ bytes/sec. When writing all 512KB in one call, you can write around $1.5 \cdot 10^{10}$, this is 10000x faster than writing one byte per call.

One of the reasons for this increase in speed is the amount of memory references involved in reading and writing the 512KB region. The methods `copy_to_user` and `copy_from_user` are responsible for reading and writing to memory in the user space. Accessing memory is a slow process, as the user virtual address has to be translated into a physical address. In the 1 bytes per write call version, there needs to be roughly $5 \cdot 10^5$ of

these references. In the 512KB per write call version, there is only one call to `copy_from_user` necessary.

Additionally, calling a system call, which generates a kernel exception, is also a slow process, as the kernel has to find the exception handler, then find the right system call to execute, and then execute the system call. The fewer system calls possible, the faster the program will run. The methods that require writing fewer bytes per system call will then be slower because more system calls are needed to write the required amount of bytes.

Source code used in part 4

I used an [EmbeTronicX tutorial](#) to understand how to use `EXPORT_SYMBOL` and compile code which relies on it. I did not directly use any of the source code from the tutorial, but I wanted to acknowledge the tutorial's help.

All other source code used in part 4 was cited in the source code section of part 3.